

Initial Steps Toward Verifying the Rust Standard Library Using Verus

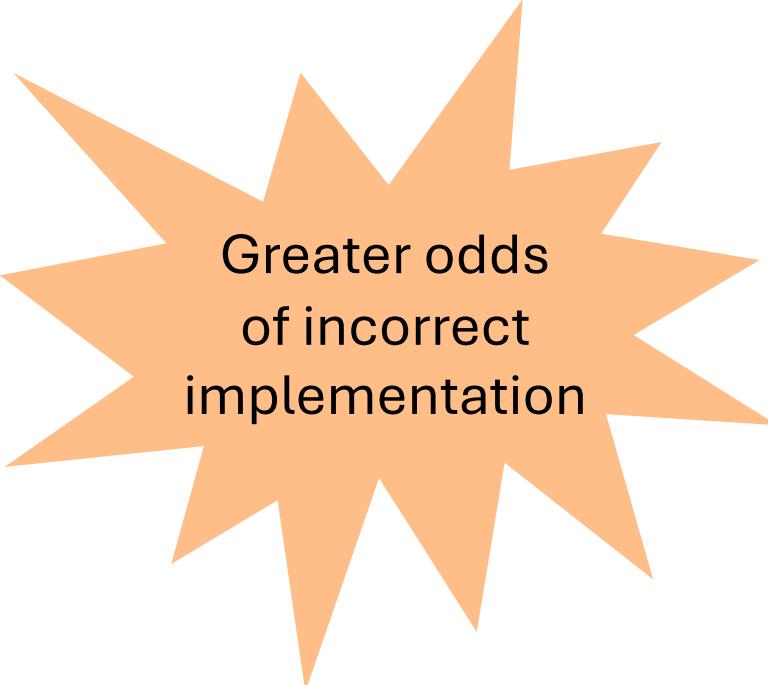
Elanor Tang, Travis Hance, and Bryan Parno

Carnegie Mellon University

Why Verify the Rust Standard Library?

- *Safe* interface
- *Unsafe* implementation

Want to provide safety and correctness guarantees



Greater odds
of incorrect
implementation



Greater
severity of
mistake



Impact extends
to most Rust
code

Safety Depends on Correctness

```
pub const fn from_utf8(v: &[u8]) -> Result<&str, Utf8Error> {
    match run_utf8_validation(v) {
        Ok(_) => {
            // SAFETY: validation succeeded.
            Ok(unsafe { from_utf8_unchecked(v) })
        }
        Err(err) => Err(err),
    }
}
```

Checks to see if a sequence of bytes is in UTF-8 format

If so, calls the unsafe function `from_utf8_unchecked()`

```
pub const unsafe fn from_utf8_unchecked(v: &[u8]) -> &str {
    // SAFETY: the caller must guarantee that the bytes `v` are valid UTF-8.
    // Also relies on `&str` and `&[u8]` having the same layout.
    unsafe { mem::transmute(v) }
}
```

Performs the unsafe operation `transmute()`

Safety Depends on Correctness

```
pub const fn from_utf8(v: &[u8]) -> Result<&str, Utf8Error> {
    match run_utf8_validation(v) {
        Ok(_) => {
            // SAFETY: validation succeeded.
            Ok(unsafe { from_utf8_unchecked(v) })
        }
        Err(err) => Err(err),
    }
}
```

Checks to see if a sequence of bytes is in UTF-8 format

If so, calls the unsafe function `from_utf8_unchecked()`

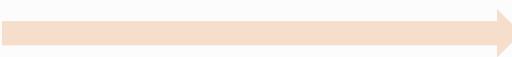
```
pub const unsafe fn from_utf8_unchecked(v: &[u8]) -> &str {
    // SAFETY: the caller must guarantee that the bytes `v` are valid UTF-8.
    // Also relies on `&str` and `&[u8]` having the same layout.
    unsafe { mem::transmute(v) }
}
```

Performs the unsafe operation `transmute()`

Safety Depends on Correctness

```
pub const fn from_utf8(v: &[u8]) -> Result<&str, Utf8Error> {  
    match run_utf8_validation(v) {  
        Ok(_) => {  
            // SAFETY: validation succeeded.  
            Ok(unsafe { from_utf8_unchecked(v) })  
        }  
        Err(err) => Err(err),  
    }  
}
```

Safety of `mem::transmute(v)` depends on correctness of `run_utf8_validation(v)`



```
pub const unsafe fn from_utf8_unchecked(v: &[u8]) -> &str {  
    // SAFETY: the caller must guarantee that the bytes `v` are valid UTF-8.  
    // Also relies on `&str` and `&[u8]` having the same layout.  
    unsafe { mem::transmute(v) }
```



Performs the unsafe operation `transmute()`

Initial Target: Verifying the String Library

Enables greater security

Presents nontrivial verification challenges

CVE-2020-36323 PUBLISHED [View JSON](#) | [User Guide](#)

CVE-2020-36317 PUBLISHED [View JSON](#) | [User Guide](#)

Re

Required CVE Record Information

CNA: MITRE Corporation

Published: 2021-04-11 Updated: 2021-04-11

Description

Challenges: Verifying the String Library

- Reasoning about pointer *provenance*

Provenance: Captures what you are allowed to do with a pointer based on the source it was derived from

Snippet from string standard library

```
unsafe {  
    let block = ptr.add(index) as *const usize;  
}
```



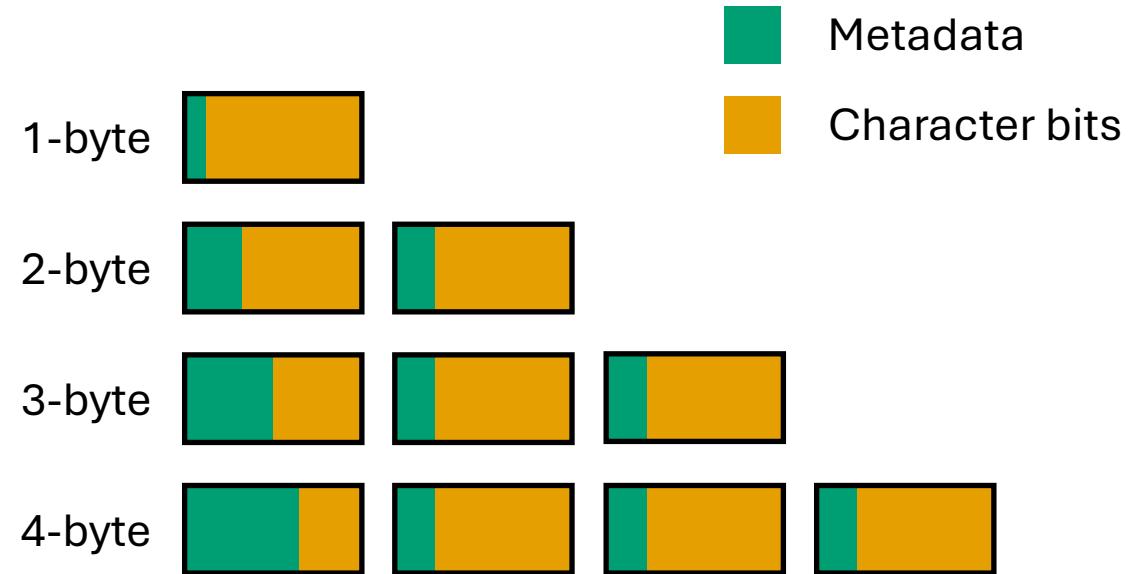
Unsafe, precondition from Rust documentation involves the provenance of `ptr`

Will talk about why it is hard later

Will talk more about `add()` later

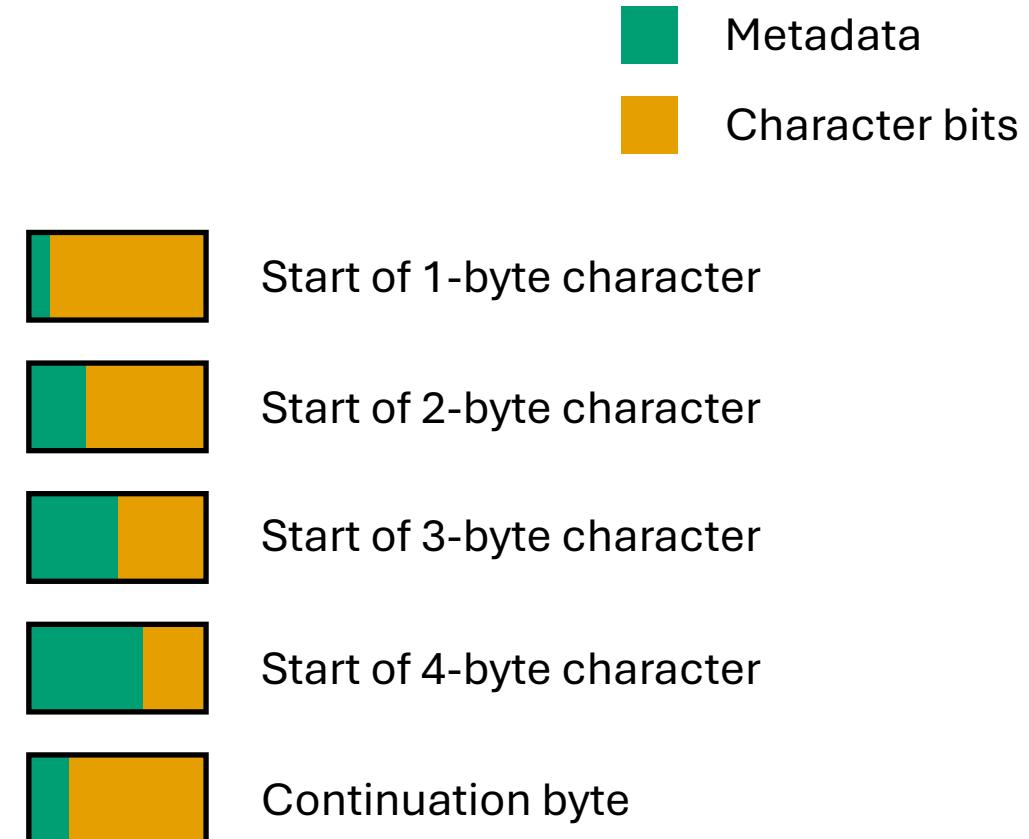
Challenges: Verifying the String Library

- Reasoning about pointer *provenance*
- Complex UTF-8 reasoning
 - Variable-width encoding of characters
 - Complicated interpretation of byte values



Challenges: Verifying the String Library

- Reasoning about pointer *provenance*
- Complex UTF-8 reasoning
 - Variable-width encoding of characters
 - Complicated interpretation of byte values
- Low-level bit manipulation

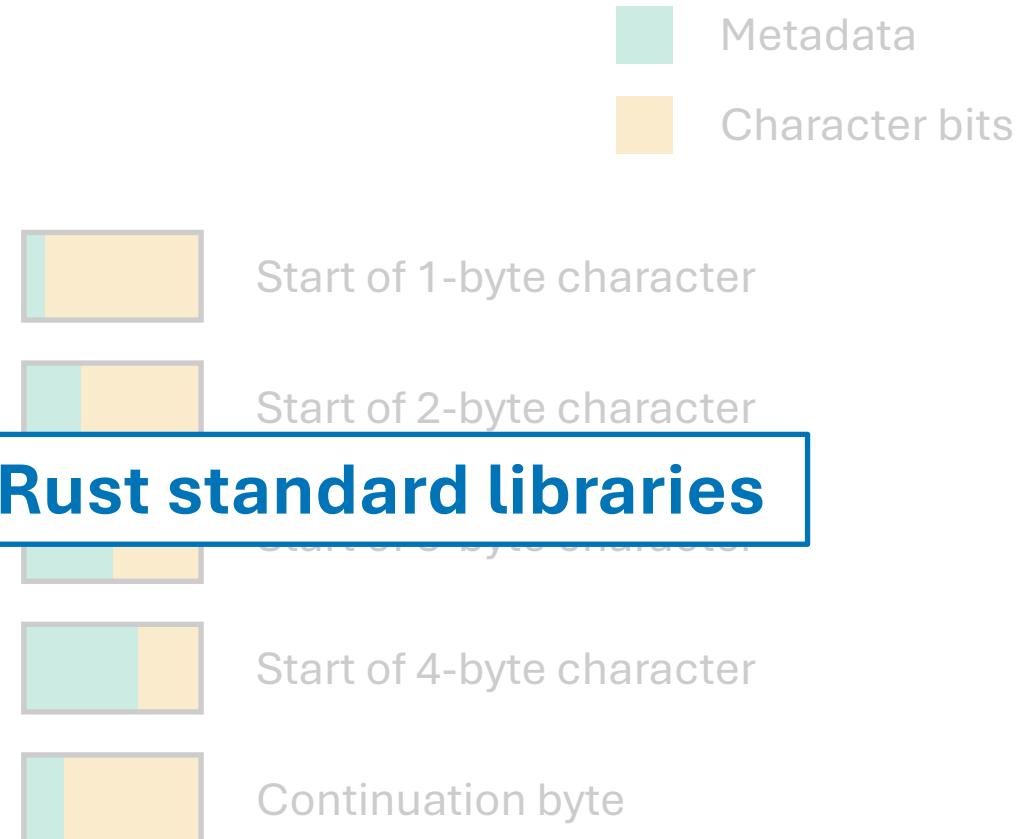


Challenges: Verifying the String Library

- Reasoning about pointer *provenance*
- Complex UTF-8 reasoning
 - Variable-width encoding of characters

Solutions apply to verifying other Rust standard libraries

- Low-level bit manipulation
 - byte values



Overview

- **Verus**
 - *Ownership ghost permissions* to reason about pointers
- **Challenges**
 - Key challenge: **Handling pointer provenance**
 - Verus challenge: Ergonomically incorporating spec/proof code into existing Rust code
- **What Went Well**
 - Successful verification of complex functions

Our Approach: Verus



- Automated SMT-based verification tool
- Uses ***ownership ghost permissions*** which are ***borrow-checked*** to safely manipulate pointers

OOSPLA'23

Verus: Verifying Rust Programs using Linear Ghost Types

ANDREA LATTUADA*, VMware Research, Switzerland
TRAVIS HANCE, Carnegie Mellon University, USA
CHANHEE CHO, Carnegie Mellon University, USA
MATTHIAS BRUN, ETH Zurich, Switzerland
ISITHA SUBASINGHE†, UNSW Sydney, Australia
YI ZHOU, Carnegie Mellon University, USA
JON HOWELL, VMware Research, USA
BRYAN PARNO, Carnegie Mellon University, USA
CHRIS HAWBLITZEL, Microsoft Research, USA

The Rust programming language provides a powerful type system that checks linearity and borrowing, allowing code to safely manipulate memory without garbage collection and making Rust ideal for developing low-level, high-assurance systems. For such systems, formal verification can be useful to prove functional correctness properties beyond type safety. This paper presents Verus, an SMT-based tool for formally verifying Rust programs. With Verus, programmers express proofs and specifications using the Rust language, allowing proofs to take advantage of Rust's linear types and borrow checking. We show how this allows proofs to manipulate linearly typed permissions that let Rust code safely manipulate memory, pointers, and concurrent resources. Verus provides a formal verification infrastructure that distinguishes between specifications which

ptr : *mut T
value: T

Track information about
what the pointer points to

Ownership Ghost Permissions

```
fn main() {  
    let (p, Tracked(mut points_to)) = allocate::<u32>(4);  
}
```

Signifies ownership
ghost permission

Stack

p = 0x1004

Heap	
Address	Value
0x1000	-
0x1004	UnInit
0x1008	-
0x100c	-

Verus' perspective

```
p = 0x1004  
points_to = {  
    ptr: 0x1004,  
    value: UnInit,  
}
```

Ownership Ghost Permissions

```
fn main() {  
    let (p, Tracked(mut points_to)) = allocate::<u32>(4);  
    ptr_mut_write(p, Tracked(&mut points_to), 5);  
}
```

Signifies ownership
ghost permission

Stack

p = 0x1004

Heap	
Address	Value
0x1000	-
0x1004	5
0x1008	-
0x100c	-

Verus' perspective

```
p = 0x1004  
points_to = {  
    ptr: 0x1004,  
    value: 5,  
}
```

(example simplified for demonstration purposes)

Features of Ownership Ghost Permissions

- Mutability of permission matches mutability of the pointer operation.
- Lifetime of permission is tied to lifetime of allocation.
- **Erase** permissions during compilation.



Enforced by ***borrow-checking***
the permissions

```
fn main() {  
    let (p, Tracked(mut points_to)) = allocate::<u32>(4);  
    ptr_mut_write(p, Tracked(&points_to), 5);    // FAILS  
    deallocate(ptr, 4, Tracked(points_to));  
    ptr_mut_write(p, Tracked(&mut points_to), 5); // FAILS  
}
```

(example simplified for demonstration purposes)

Verus: Benefits and Limitations



- ✓ Can reason about safe and unsafe code
- ✓ Has the automation to scale (~250K total Verus LOC, codebases up to 60K LOC)
- ✗ Needs more features to handle pointer provenance reasoning

SOSP'24

Check for Updates

Verus: A Practical Foundation for Systems Verification

Andrea Lattuada*
MPI-SWS

Matthias Brun
ETH Zurich

Pranav Srinivasan
University of Michigan

Chris Hawblitzel
Microsoft Research

Oded Padon*
Weizmann Institute of Science

Travis Hance
Carnegie Mellon University

Chanhee Cho
Carnegie Mellon University

Reto Achermann
University of British Columbia

Jon Howell
VMware Research

Bryan Parno
Carnegie Mellon University

Jay Bosamiya†
Microsoft Research

Hayley LeBlanc
University of Texas at Austin

Tej Chajed
University of Wisconsin-Madison

Jacob R. Lorch
Microsoft Research

Artifacts Evaluated Functional V1.1

Artifacts Available V1.1

Results Reproduced V1.1

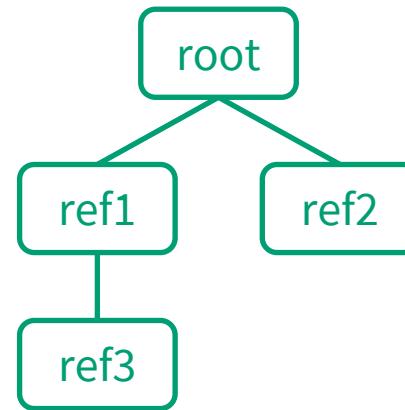
CCS Concepts: • Software and its engineering → Formal software verification.

Key Challenge: Handling Provenance

- Rust has no formal provenance model
- Many provenance models are complicated
 - E.g., tree and stacked borrows

“The exact structure of provenance is not yet specified, but the permission defined by a pointer’s provenance has a *spatial* component, a *temporal* component, and a *mutability* component.”

—Rust pointer documentation



Our Solution: Handling Provenance

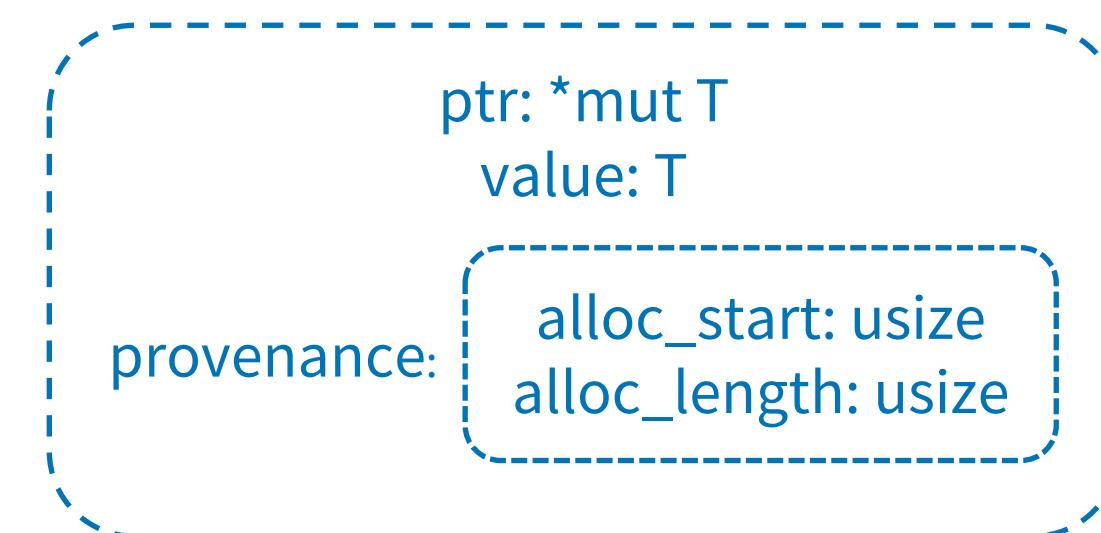
Key idea: Abstract over the provenance model

✓ Spatial → Extend ownership ghost permissions
with provenance information

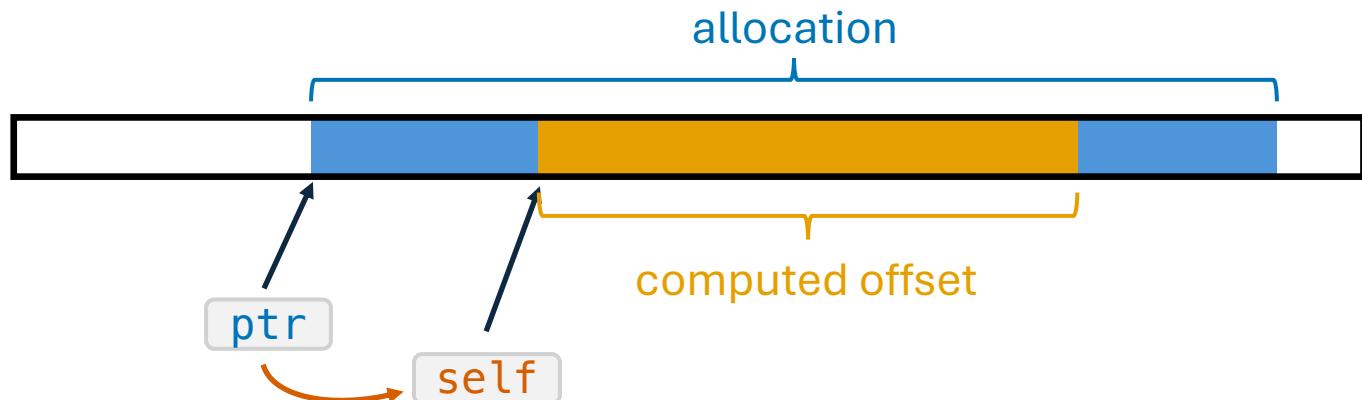
✓ Temporal

✓ Mutability

Taken care of by Rust's
borrow-checker on the
lifetime and mutability
of ghost permissions



Example: `ptr.add()`



If the computed offset is non-zero, then

- 1 `self` must be derived from some pointer to an allocated object
- 2 Memory range between `self` and the result must be within bounds

Advance `self: *const T` by `count` elements of type `T`

```
impl<T: Sized> *const T {  
    pub const unsafe fn add(self, count: usize) -> output  
}
```

{...}

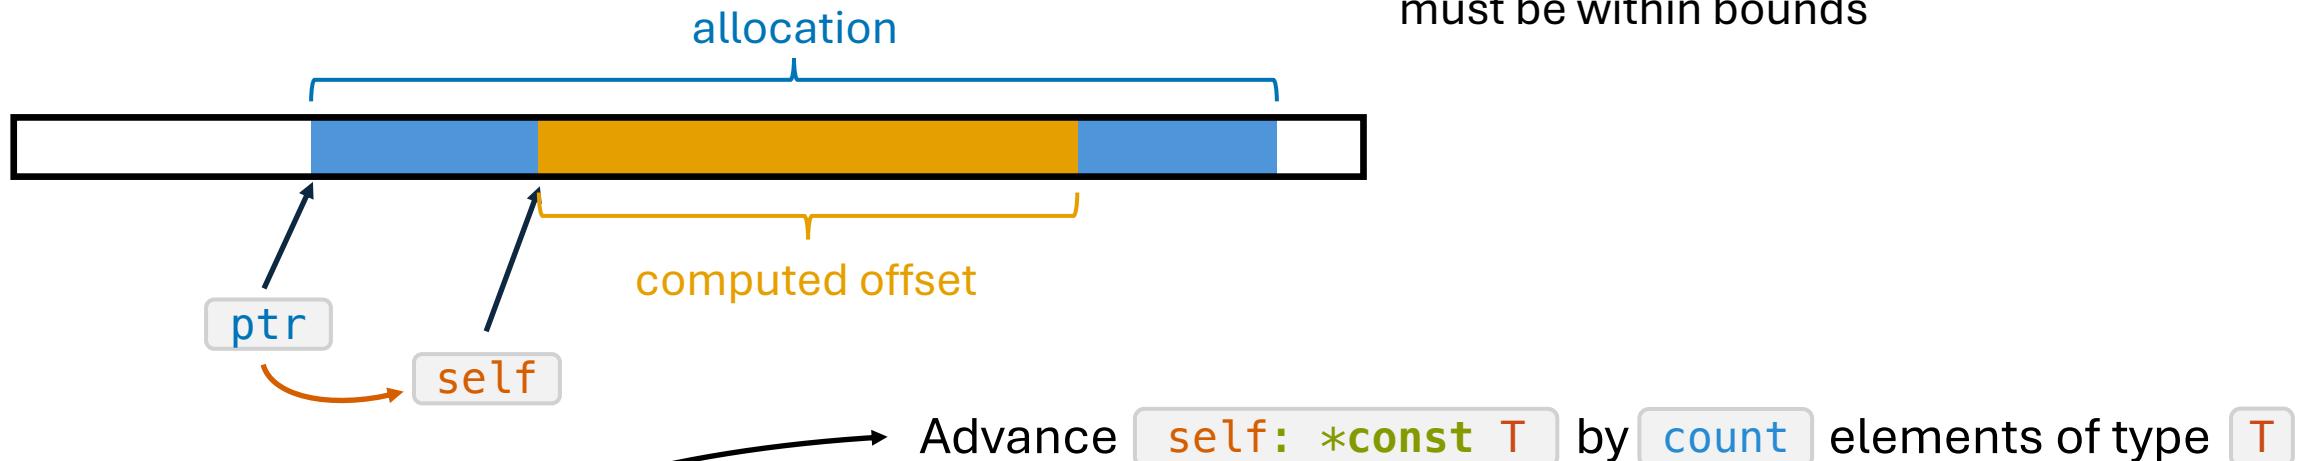
}

Example: `ptr.add()`



```
impl<T: Sized> *const T {  
    pub const unsafe fn add_verus(self, count: usize, Tracked(perm): Tracked<&PointsToRaw>) -> (output: Self)  
    requires  
        { ... }  
        {  
            provenance:  
                alloc_start: usize  
                alloc_length: usize  
        }  
}
```

Example: `ptr.add()`



```
impl<T: Sized> *const T {  
    pub const unsafe fn add_verus(self, count: usize, Tracked(perm): Tracked<&PointsToRaw>) -> (output: Self)  
    requires  
        count * size_of::<T>() > 0  
    ==>  
    {...}  
}
```

provenance: `alloc_start: usize`
`alloc_length: usize`

Example: `ptr.add()`

If the computed offset is non-zero, then

- ✓ 1 `self` must be derived from some pointer to an allocated object
- 2 Memory range between `self` and the result must be within bounds



```
impl<T: Sized> *const T {  
    pub const unsafe fn add_verus(self, count: usize, Tracked(perm): Tracked<&PointsToRaw>) -> (output: Self)  
    requires  
        count * size_of::<T>() > 0  
    ==> perm.provenance() == self@.provenance}
```

provenance: {
 alloc_start: usize
 alloc_length: usize

{...}
}

Example: `ptr.add()`

If the computed offset is non-zero, then

- ✓ 1 `self` must be derived from some pointer to an allocated object
- ✓ 2 Memory range between `self` and the result must be within bounds



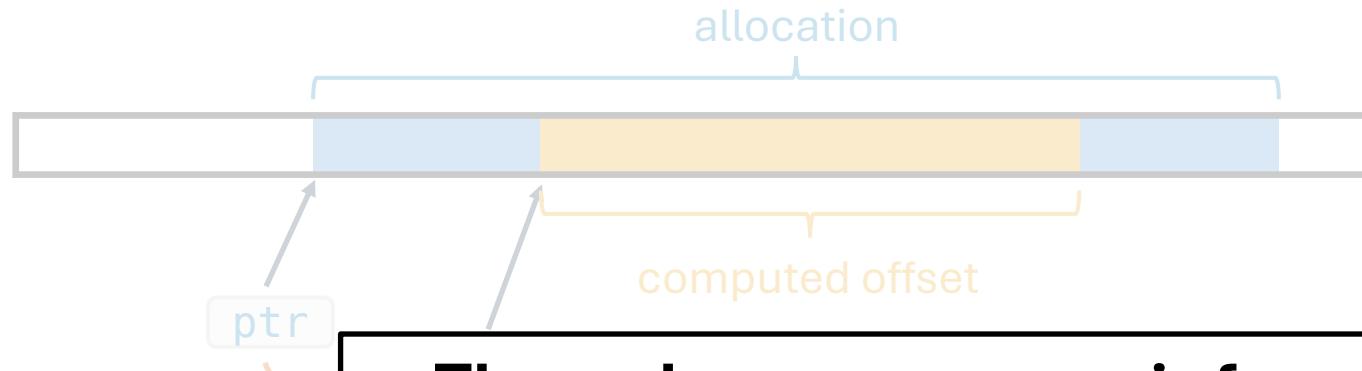
```
impl<T: Sized> *const T {  
    pub const unsafe fn add_verus(self, count: usize, Tracked(perm): Tracked<&PointsToRaw>) -> (output: Self)  
    requires  
        count * size_of::<T>() > 0  
        ==> perm.provenance() == self@.provenance  
        && (self@.addr + count * size_of::<T>())  
            < self@.provenance.alloc_start() + self@.provenance.alloc_len()  
        && self@.addr >= self@.provenance.alloc_start(),  
    {...}  
}
```

provenance: {
 alloc_start: usize
 alloc_length: usize
}

Example: `ptr.add()`

If the computed offset is non-zero, then

- ✓ 1 `self` must be derived from some pointer to an allocated object
- ✓ 2 Memory range between `self` and the result must be within bounds



The only provenance information we needed was `alloc_start` and `alloc_length`

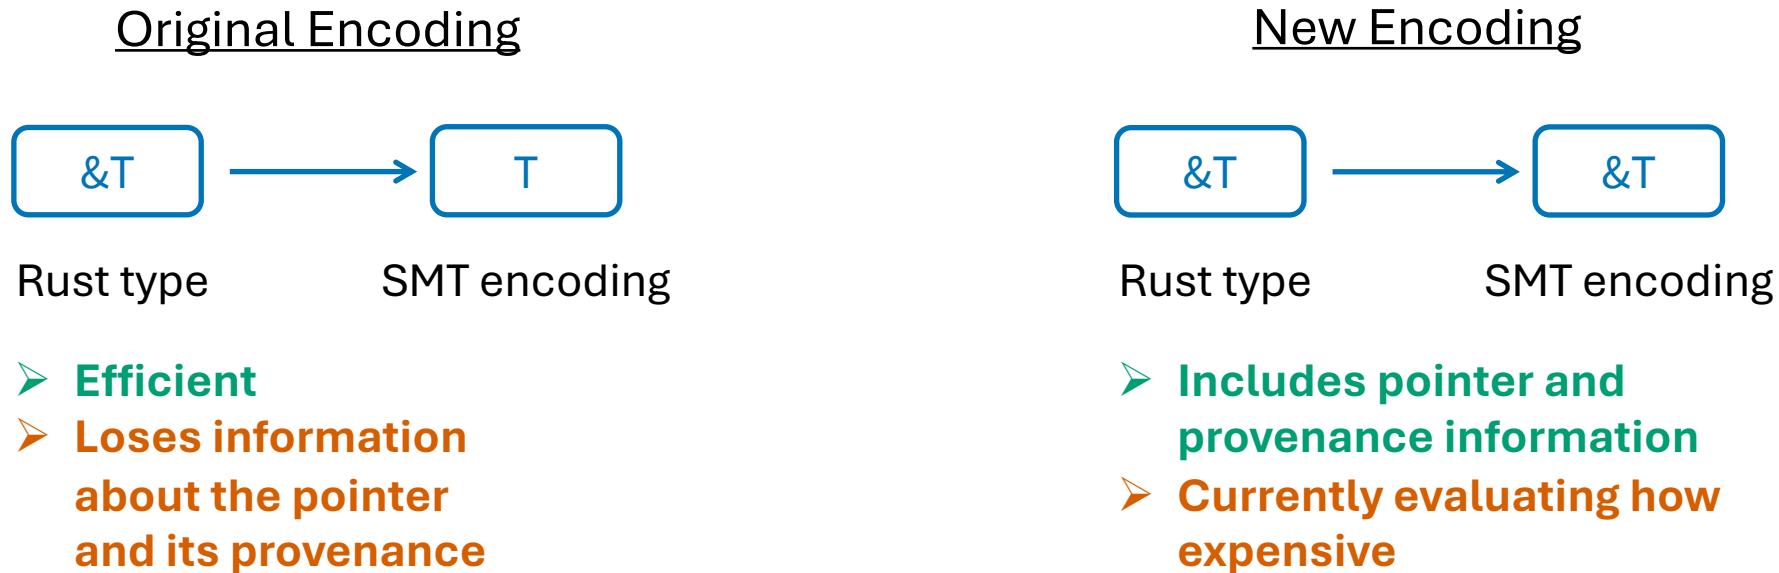
```
impl<T: Sized> *const T<T> {  
    pub const unsafe fn add_verus(self, count: usize, Tracked(perm): Tracked<&PointsToRaw>) -> (output: Self)  
    requires  
        count * size_of::<T>() > 0  
        ==> perm.provenance() == self@.provenance  
        && (self@.addr + count * size_of::<T>())  
            < self@.provenance.alloc_start() + self@.provenance.alloc_len()  
        && self@.addr >= self@.provenance.alloc_start(),  
    {...}  
}
```

s of type `T`

provenance:
 `alloc_start: usize`
 `alloc_length: usize`

Additional Wrinkle

Carrying around pointer/provenance information at the SMT level may get expensive.



Potential challenge: Incorporating both encodings

Verus Challenge: Ergonomically Incorporating Spec/Proof Code into Existing Rust Code

- Need to use Verus versions of functions

`ptr.add(count)`



`ptr.add_verus(count), Tracked(&perm)`

`*block`



`*ptr_ref(block, Tracked(&perm))`

`*ptr = 5`



`ptr_mut_write(p, Tracked(&mut perm), 5)`

- Rust currently lacks support for ghost code

Proposed Solution: Ergonomically Incorporating Spec/Proof Code into Existing Rust Code

```
pub fn ptr_mut_write<T>(ptr: *mut T, Tracked(points_to): Tracked<&mut PointsTo<T>>, v: T)
```

- Proposal to move ownership ghost function parameters out of the main function signature

```
pub fn ptr_mut_write<T>(ptr: *mut T, v: T)  
    with_ghost_arg (points_to: &mut PointsTo<T>)
```

- Support in progress for attribute-based syntax

```
#[with_ghost_arg(points_to: &mut PointsTo<T>)]  
pub fn ptr_mut_write<T>(ptr: *mut T, v: T)
```

Proposed Solution: Ergonomically Incorporating Spec/Proof Code into Existing Rust Code

```
pub fn ptr_mut_write<T>(ptr: *mut T, Tracked(points_to): Tracked<&mut PointsTo<T>>, v: T)
```

- Proposal to move ownership ghost function parameters out of the main function signature

**Need to do this in a way that still enables type-checking,
so we can keep borrow-checking our permissions**

- Support in progress for attribute-based syntax

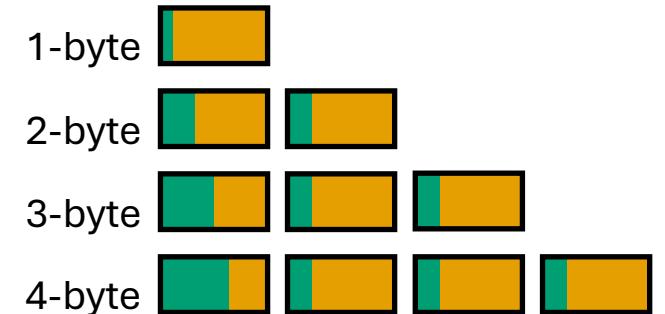
```
#[with_ghost_arg(points_to: &mut PointsTo<T>)]
pub fn ptr_mut_write<T>(ptr: *mut T, v: T)
```

Putting It Together: What Went Well

Metadata
Character bits

Successful verification of complex functions

Clean mathematical spec for UTF-8



```
pub const fn run_utf8_validation(v_ref: &[u8]) -> (result: Result<(), Utf8Error>)
ensures
    is_ok(&result) <==> valid_utf8(v_ref@),
{
```

Low-level *automatic* bitwise proofs

}

Low-level bitwise proofs

All assertions about bitwise operations were **automatically** proved

- Context of example
 - Showing block of bytes is ASCII
 - Endianness reasoning

$$\begin{array}{c} \text{x_high} \quad \text{x_low} \\ \text{&} \\ \text{y_high} \quad \text{y_low} \end{array} == 0$$

Automatic

$$\begin{array}{c} \text{x_high} \quad \text{---} \\ \text{&} \\ \text{y_high} \quad \text{---} \\ \text{&&} \\ \text{---} \quad \text{x_low} \\ \text{&} \\ \text{---} \quad \text{y_low} \end{array} == 0$$

Putting It Together: What Went Well

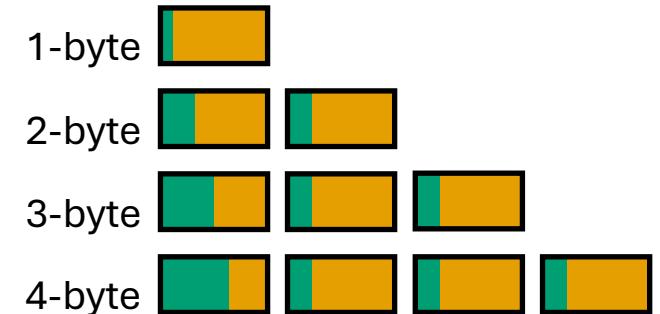
Successful verification of complex functions

Clean mathematical spec for UTF-8

```
pub const fn run_utf8_validation(v_ref: &[u8]) -> (result: Result<(), Utf8Error>)
ensures
    is_ok(&result) <==> valid_utf8(v_ref@),
{
```

Low-level *automatic* bitwise proofs

Pointer *provenance* reasoning



Pointer Provenance Reasoning via Ownership Ghost Permissions

ptr: *const u8

```
let block = ptr.add(index) as *const usize;  
let zu = contains_nonascii(*block);
```

Unverified code in run_utf8_validation()

```
let block = ptr.add_verus(index, Tracked(perm.into_raw_shared())) as *const usize;  
let tracked zu_sub_perm = perm.subrange(index, size_of::<usize>());  
proof {  
    lemma_mod_eqv_0_trans(zu_sub_perm.ptr()@.addr as int,  
                          size_of::<usize>() as int, align_of::<usize>() as int);  
}  
let tracked zu_usize_perm = zu_sub_perm.cast_points_to::<usize>();  
let x = *ptr_ref(block, Tracked(zu_usize_perm));  
let zu = contains_nonascii(x);
```

Verified code in run_utf8_validation()

Pointer Provenance Reasoning via Ownership Ghost Permissions

ptr: *const u8
unsafe

```
let block = ptr.add(index) as *const usize;  
let zu = contains_nonascii(*block);
```

Unverified code in run_utf8_validation()

```
let block = ptr.add verus(index, Tracked(perm.into raw shared())) as *const usize;  
let tracked zu_sub_perm = perm.subrange(index, size_of::<usize>());  
proof {  
    lemma_mod_eqv_0_trans(zu_sub_perm.ptr()@.addr as int,  
                          size_of::<usize>() as int, align_of::<usize>() as int);  
}  
let tracked zu_usize_perm = zu_sub_perm.cast_points_to::<usize>();  
let x = *ptr_ref(block, Tracked(zu_usize_perm));  
let zu = contains_nonascii(x);
```

Verified code in run_utf8_validation()

Pointer Provenance Reasoning via Ownership Ghost Permissions

ptr: *const u8

```
let block = ptr.add(index) as *const usize;  
let zu = contains_nonascii(*block);
```

Unverified code in run_utf8_validation()

```
let block = ptr.add_verus(index, Tracked(perm.into_raw_shared())) as *const usize;  
let tracked zu_sub_perm = perm.subrange(index, size_of::<usize>());  
proof {  
    lemma_mod_eqv_0_trans(zu_sub_perm.ptr()@.addr as int,  
                          size_of::<usize>() as int, align_of::<usize>() as int);  
}  
let tracked zu_usize_perm = zu_sub_perm.cast_points_to::<usize>();  
let x = *ptr_ref(block, Tracked(zu_usize_perm));  
let zu = contains_nonascii(x);
```

Verified code in run_utf8_validation()

Recap: Essential Verification Tool Features

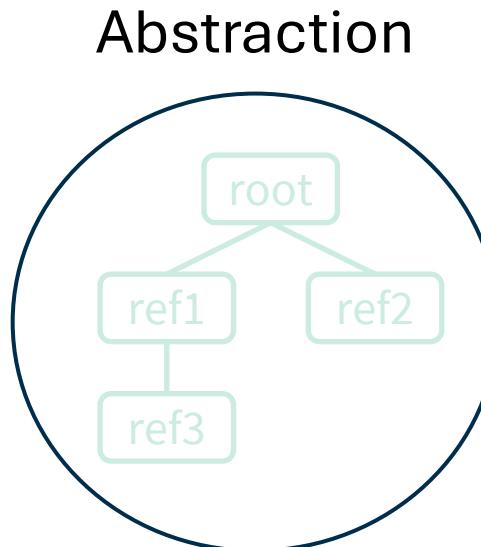
- ✓ Mathematical specification language
- ✓ Automation
- ✓ Bit-level reasoning
 - Ergonomically integrate spec/proof code into existing Rust code
 - Provenance-specific pointer reasoning



- Proposal to decouple executable and ownership ghost function parameters
 - Support in progress for attribute-based syntax
-
- ✓ Verus has pointer reasoning via ownership ghost permissions
 - ✓ Added provenance information to permissions
 - Next step: experiment with more pointer manipulations

Recap: Key Discussion Question

How to add abstraction over the provenance model?



We would love to hear your thoughts on the best approach!

Thank you!